Do LLMs know when they can coordinate? Confidence calibration on a graded stimulus battery

6 stimuli (easy to hard) x 2 models, one−shot dyadic coordination, position−counterbalanced. Points above the dashed line = overconfident.
Top: confidence vs. true match likelihood. Bottom: willingness to bet on coordination as the safe payoff rises.

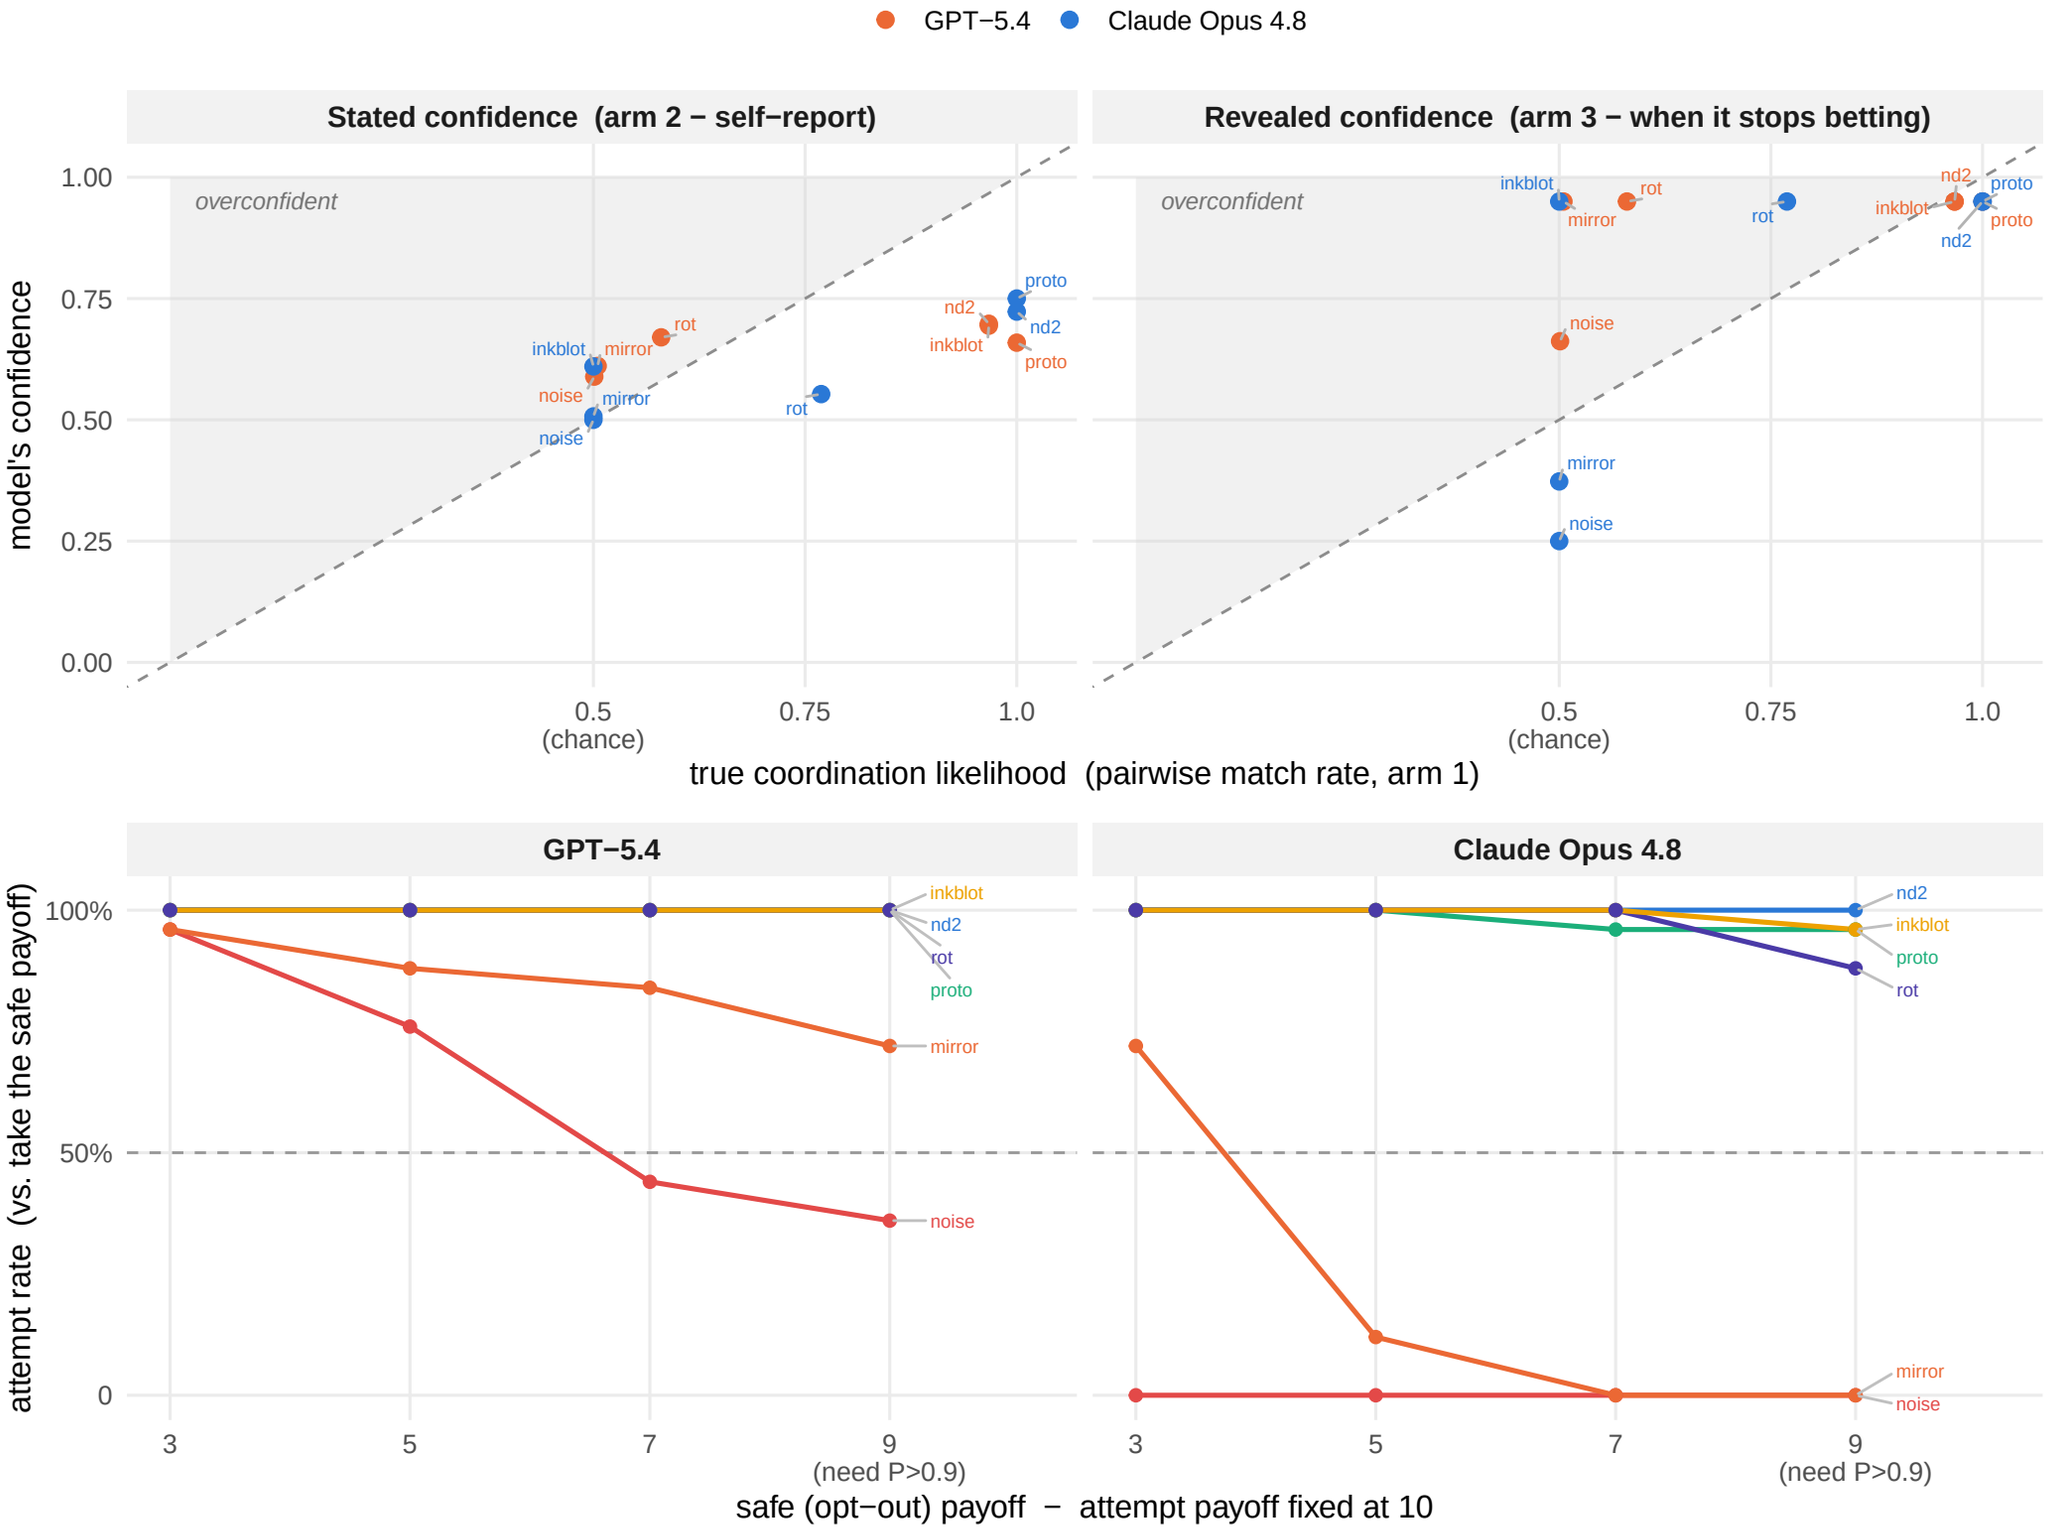

Arm 1 choice (N=60/cell) −> true pairwise match. Arm 2 stated P (N=50). Arm 3 attempt−vs−opt−out sweep (N=25/cell x safe in {3,5,7,9}). 2,520 calls. gpt−5.4 and claude−opus−4−8 at low reasoning effort.