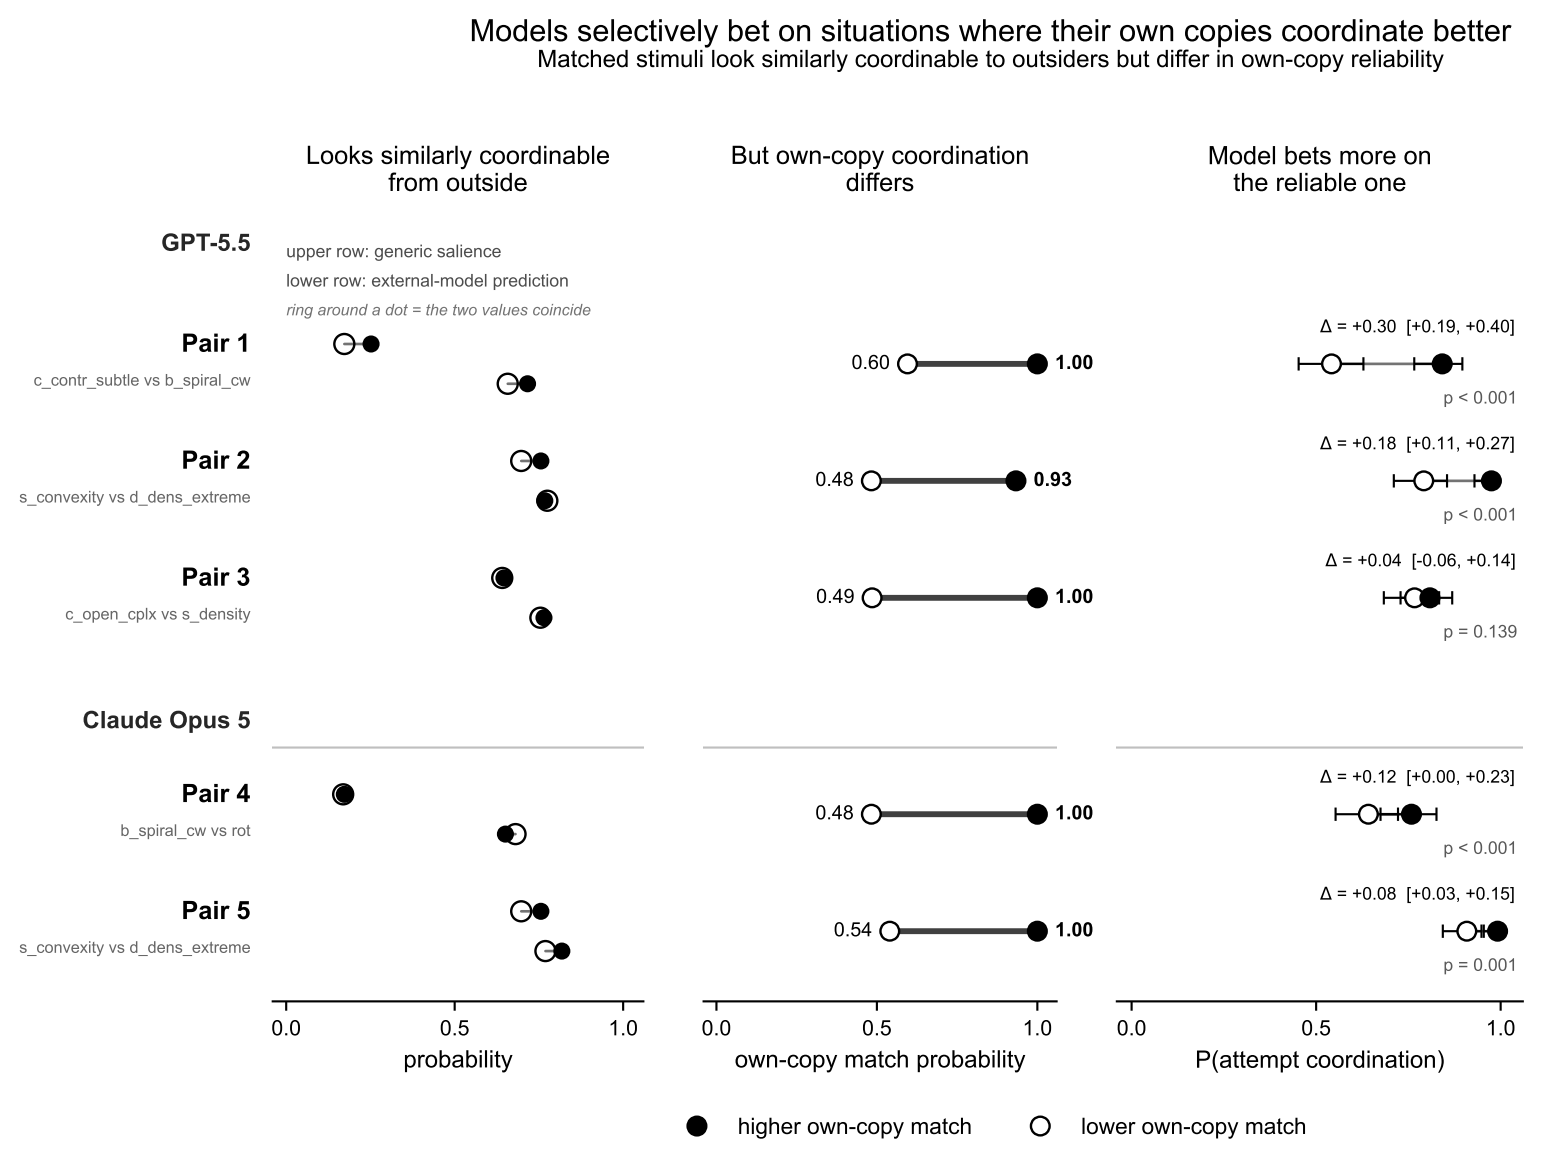
Models selectively bet on situations where their own copies coordinate better
Matched stimuli look similarly coordinable to outsiders but differ in own-copy reliability
Looks similarly coordinable from outside
But own-copy coordination differs
Model bets more on the reliable one
GPT-5.5
upper row: generic salience
lower row: external-model prediction
ring around a dot = the two values coincide
Pair 1
c_contr_subtle vs b_spiral_cw
0.60
1.00
Δ = +0.30 [+0.19, +0.40]
p < 0.001
Pair 2
s_convexity vs d_dens_extreme
0.48
0.93
Δ = +0.18 [+0.11, +0.27]
p < 0.001
Pair 3
c_open_cplx vs s_density
0.49
1.00
Δ = +0.04 [-0.06, +0.14]
p = 0.139
Claude Opus 5
Pair 4
b_spiral_cw vs rot
0.48
1.00
Δ = +0.12 [+0.00, +0.23]
p < 0.001
Pair 5
s_convexity vs d_dens_extreme
0.54
1.00
Δ = +0.08 [+0.03, +0.15]
p = 0.001
0.0
0.5
1.0
probability
own-copy match probability
P(attempt coordination)
higher own-copy match
lower own-copy match